

Logistic Regression Stochastic Gradient Descent

March 24, 2021

1 Implement SGD Classifier with Logloss and L2 regularization Using SGD without using sklearn

There will be some functions that start with the word “grader” ex: grader_weights(), grader_sigmoid(), grader_logloss() etc, you should not change those function definition. Every Grader function has to return True.

Importing packages

```
[51]: import numpy as np
import pandas as pd
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn import linear_model
```

Creating custom dataset

```
[52]: # please don't change random_state
X, y = make_classification(n_samples=50000, n_features=15, n_informative=10,
    ↳ n_redundant=5,
                                n_classes=2, weights=[0.7], class_sep=0.7,
    ↳ random_state=15)
# make_classification is used to create custom dataset
# Please check this link (https://scikit-learn.org/stable/modules/generated/
    ↳ sklearn.datasets.make_classification.html) for more details
```

```
[53]: X.shape, y.shape
```

```
[53]: ((50000, 15), (50000,))
```

Splitting data into train and test

```
[54]: # please don't change random_state
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
    ↳ random_state=15)
```

```
[55]: # Standardizing the data.
scaler = StandardScaler()
```

```
x_train = scaler.fit_transform(X_train)
x_test = scaler.transform(X_test)
```

```
[56]: X_train.shape, y_train.shape, X_test.shape, y_test.shape
```

```
[56]: ((37500, 15), (37500,), (12500, 15), (12500,))
```

2 SGD classifier

```
[57]: # alpha : float
      # Constant that multiplies the regularization term.

      # eta0 : double
      # The initial learning rate for the 'constant', 'invscaling' or 'adaptive'
      ↪ schedules.

      clf = linear_model.SGDClassifier(eta0=0.0001, alpha=0.0001, loss='log',
      ↪ random_state=15, penalty='l2', tol=1e-3, verbose=2, learning_rate='constant')
      clf
      # Please check this documentation (https://scikit-learn.org/stable/modules/
      ↪ generated/sklearn.linear_model.SGDClassifier.html)
```

```
[57]: SGDClassifier(eta0=0.0001, learning_rate='constant', loss='log',
                  random_state=15, verbose=2)
```

```
[58]: clf.fit(X=X_train, y=y_train) # fitting our model
```

```
-- Epoch 1
Norm: 0.77, NNZs: 15, Bias: -0.316653, T: 37500, Avg. loss: 0.455552
Total training time: 0.01 seconds.
-- Epoch 2
Norm: 0.91, NNZs: 15, Bias: -0.472747, T: 75000, Avg. loss: 0.394686
Total training time: 0.01 seconds.
-- Epoch 3
Norm: 0.98, NNZs: 15, Bias: -0.580082, T: 112500, Avg. loss: 0.385711
Total training time: 0.01 seconds.
-- Epoch 4
Norm: 1.02, NNZs: 15, Bias: -0.658292, T: 150000, Avg. loss: 0.382083
Total training time: 0.02 seconds.
-- Epoch 5
Norm: 1.04, NNZs: 15, Bias: -0.719528, T: 187500, Avg. loss: 0.380486
Total training time: 0.02 seconds.
-- Epoch 6
Norm: 1.05, NNZs: 15, Bias: -0.763409, T: 225000, Avg. loss: 0.379578
Total training time: 0.03 seconds.
-- Epoch 7
Norm: 1.06, NNZs: 15, Bias: -0.795106, T: 262500, Avg. loss: 0.379150
```

```
Total training time: 0.03 seconds.
-- Epoch 8
Norm: 1.06, NNZs: 15, Bias: -0.819925, T: 300000, Avg. loss: 0.378856
Total training time: 0.04 seconds.
-- Epoch 9
Norm: 1.07, NNZs: 15, Bias: -0.837805, T: 337500, Avg. loss: 0.378585
Total training time: 0.04 seconds.
-- Epoch 10
Norm: 1.08, NNZs: 15, Bias: -0.853138, T: 375000, Avg. loss: 0.378630
Total training time: 0.05 seconds.
Convergence after 10 epochs took 0.05 seconds
```

```
[58]: SGDClassifier(eta0=0.0001, learning_rate='constant', loss='log',
                    random_state=15, verbose=2)
```

```
[59]: clf.coef_, clf.coef_.shape, clf.intercept_
      #clf.coef_ will return the weights
      #clf.coef_.shape will return the shape of weights
      #clf.intercept_ will return the intercept term
```

```
[59]: (array([[ -0.42336692,  0.18547565, -0.14859036,  0.34144407, -0.2081867 ,
                0.56016579, -0.45242483, -0.09408813,  0.2092732 ,  0.18084126,
                0.19705191,  0.00421916, -0.0796037 ,  0.33852802,  0.02266721]]),
      (1, 15),
      array([-0.8531383]))
```

```
# This is formatted as code
```

2.1 Implement Logistic Regression with L2 regularization Using SGD: without using sklearn

1. We will be giving you some functions, please write code in that functions only.
2. After every function, we will be giving you expected output, please make sure that you get that output.
 - Initialize the weight_vector and intercept term to zeros (Write your code in def initialize_weights())
 - Create a loss function (Write your code in def logloss())

$logloss = -1 * \frac{1}{n} \sum_{foreach Y_t, Y_{pred}} (Y_t \log_{10}(Y_{pred}) + (1 - Y_t) \log_{10}(1 - Y_{pred}))$ - for each epoch:

- for each batch of data points in train: (keep batch size=1)

- calculate the gradient of loss function w.r.t each weight in weight vector (write your code in def grad_weight())

$\frac{dw^{(t)}}{dt} = x_n(y_n - ((w^{(t)})^T x_n + b^{(t)})) - \frac{1}{N} w^{(t)}$

- Calculate the gradient of the intercept (write your code in def grad_intercept())

Grader function - 2

```
[64]: def grader_sigmoid(z):  
    val=sigmoid(z)  
    assert(val==0.8807970779778823)  
    return True  
grader_sigmoid(2)
```

[64]: True

Compute loss

$$\text{logloss} = -1 * \frac{1}{n} \sum_{\text{foreach } Y_t, Y_{\text{pred}}} (Y_t \log_{10}(Y_{\text{pred}}) + (1 - Y_t) \log_{10}(1 - Y_{\text{pred}}))$$

```
[65]: def logloss(y_true,y_pred):  
    '''In this function, we will compute log loss '''  
    loss = 0  
    n = len(y_true)  
    for i in range(n):  
        loss += ( ( np.dot(y_true[i], np.log10(y_pred[i])) ) + ( np.  
→dot((1-y_true[i]),np.log10(1-y_pred[i])) ) )  
    loss = -1/n * loss  
    return loss
```

Grader function - 3

```
[66]: def grader_logloss(true,pred):  
    loss=logloss(true,pred)  
    assert(loss==0.07644900402910389)  
    return True  
true=[1,1,0,1,0]  
pred=[0.9,0.8,0.1,0.8,0.2]  
grader_logloss(true,pred)
```

[66]: True

Compute gradient w.r.to 'w'

$$dw^{(t)} = x_n(y_n - ((w^{(t)})^T x_n + b^t)) - \frac{1}{N} w^{(t)}$$

```
[67]: def gradient_dw(x,y,w,b,alpha,N):  
    '''In this function, we will compute the gradient w.r.to w '''  
    dw = (np.dot(x,(y - sigmoid(np.dot(w,x) + b))) - (alpha/N)*w)  
    return dw
```

Grader function - 4

```
[68]: def grader_dw(x,y,w,b,alpha,N):  
    grad_dw=gradient_dw(x,y,w,b,alpha,N)  
    assert(np.sum(grad_dw)==2.613689585)  
    return True
```

```
grad_x=np.array([-2.07864835,  3.31604252, -0.79104357, -3.87045546, -1.
↪14783286,
                -2.81434437, -0.86771071, -0.04073287,  0.84827878,  1.99451725,
                3.67152472,  0.01451875,  2.01062888,  0.07373904, -5.54586092])
grad_y=0
grad_w,grad_b=initialize_weights(grad_x)
alpha=0.0001
N=len(X_train)
grader_dw(grad_x,grad_y,grad_w,grad_b,alpha,N)
```

[68]: True

Compute gradient w.r.to 'b'

$$\delta b^{(t)} = y_n - ((w^{(t)})^T x_n + b^{(t)})$$

```
[69]: def gradient_db(x,y,w,b):
        '''In this function, we will compute gradient w.r.to b '''
        db = y - sigmoid(np.dot(w,x) + b)
        return db
```

Grader function - 5

```
[70]: def grader_db(x,y,w,b):
        grad_db=gradient_db(x,y,w,b)
        assert(grad_db==0.5)
        return True
grad_x=np.array([-2.07864835,  3.31604252, -0.79104357, -3.87045546, -1.
↪14783286,
                -2.81434437, -0.86771071, -0.04073287,  0.84827878,  1.99451725,
                3.67152472,  0.01451875,  2.01062888,  0.07373904, -5.54586092])
grad_y=0
grad_w,grad_b=initialize_weights(grad_x)
alpha=0.0001
N=len(X_train)
grader_db(grad_x,grad_y,grad_w,grad_b)
```

[70]: True

Implementing logistic regression

```
[71]: def train(X_train,y_train,X_test,y_test,epochs,alpha,eta0):
        ''' In this function, we will implement logistic regression'''
        #Here eta0 is learning rate
        #implement the code as follows
        # initialize the weights (call the initialize_weights(X_train[0]) function)
        # for every epoch
            # for every data point(X_train,y_train)
                #compute gradient w.r.to w (call the gradient_dw() function)
```

```

        #compute gradient w.r.to b (call the gradient_db() function)
        #update w, b
        # predict the output of x_train[for all data points in X_train] using
↪w, b
        #compute the loss between predicted and actual values (call the loss
↪function)
        # store all the train loss values in a list
        # predict the output of x_test[for all data points in X_test] using w,b
        #compute the loss between predicted and actual values (call the loss
↪function)
        # store all the test loss values in a list
        # you can also compare previous loss and current loss, if loss is not
↪updating then stop the process and return w,b
        N = len(X_train)
        w, b = initialize_weights(X_train[0])

        # x_train_loss and x_test_loss will store the loss values at the end of
↪each epoch
        x_train_loss = []
        x_test_loss = []

        # prev_train_loss and prev_test_loss will keep track of the loss values in
↪the previous epoch
        prev_train_loss = 0
        prev_test_loss = 0

        # epochs_ran will be incremented each time an epoch is finished
        epochs_ran = 0

        # iterating over the number of epochs
        for i in range(epochs):
            for j in range(len(X_train)):
                # calculating the dw and db values and updating w and b values
                dw = gradient_dw(X_train[j], y_train[j], w, b, alpha, N)
                db = gradient_db(X_train[j], y_train[j], w, b)
                w = w + (eta0 * dw)
                b = b + (eta0 * db)

                # x_train_pred and x_test_pred will store the probabilities obtained
↪using this epoch's w and b values
                x_train_pred = []
                x_test_pred = []

                # calculating and storing x_train_pred and x_test_pred values
                for k in range(len(X_train)):
                    x_train_pred.append(sigmoid(np.dot(w, X_train[k])+b))

```

```

for k in range(len(X_test)):
    x_test_pred.append(sigmoid(np.dot(w, X_test[k])+b))

    # calculating the logloss values of train and test probabilities
    →obtained
    current_train_loss = logloss(y_train, x_train_pred)
    current_test_loss = logloss(y_test, x_test_pred)

    # storing the logloss values obtained
    x_train_loss.append(current_train_loss)
    x_test_loss.append(current_test_loss)

    # printing the informations of this epoch
    print(f"Epoch: {i}\nTrain Loss: {current_train_loss}\nTest Loss:
    →{current_test_loss}\nPrev Train Loss: {prev_train_loss}\nPrev Test Loss:
    →{prev_test_loss}\n", "="*10, "\n")

    # if this is the first epoch, storing the current loss values in the
    →prev_train_loss and prev_test_loss values
    if(i == 0):
        prev_train_loss = current_train_loss
        prev_test_loss = current_test_loss
    # if this is not the first epoch,
    else:
        # checking if the difference of train loss between current epoch
        →and previous epoch is less than  $10^{-5}$ 
        if(prev_train_loss-current_train_loss < 0.00001):
            # if it is, breaking the execution
            print("No improvement in train loss")
            break
        else:
            # if it is not, then storing the current train and test losses
            →as previous loss values and moving on to the next epoch
            prev_train_loss = current_train_loss
            prev_test_loss = current_test_loss

    # incrementing the epochs_ran value
    epochs_ran += 1
    # returning the values
    return w,b,x_train_loss,x_test_loss, epochs_ran

```

```

[72]: alpha=0.0001
      eta0=0.0001
      N=len(X_train)
      epochs=50

```



```
w,b,x_train_loss,x_test_loss,epochs_ran =  
↳train(X_train,y_train,X_test,y_test,epochs,alpha,eta0)
```

```
Epoch: 0  
Train Loss: 0.17545748442854608  
Test Loss: 0.1759547442321374  
Prev Train Loss: 0  
Prev Test Loss: 0  
=====
```

```
Epoch: 1  
Train Loss: 0.16867157050333045  
Test Loss: 0.16939931358951013  
Prev Train Loss: 0.17545748442854608  
Prev Test Loss: 0.1759547442321374  
=====
```

```
Epoch: 2  
Train Loss: 0.1663916799246292  
Test Loss: 0.16720591194885745  
Prev Train Loss: 0.16867157050333045  
Prev Test Loss: 0.16939931358951013  
=====
```

```
Epoch: 3  
Train Loss: 0.16536827537403162  
Test Loss: 0.16621717799334956  
Prev Train Loss: 0.1663916799246292  
Prev Test Loss: 0.16720591194885745  
=====
```

```
Epoch: 4  
Train Loss: 0.16485707459547086  
Test Loss: 0.16571959463978406  
Prev Train Loss: 0.16536827537403162  
Prev Test Loss: 0.16621717799334956  
=====
```

```
Epoch: 5  
Train Loss: 0.16458820012928269  
Test Loss: 0.16545557095508645  
Prev Train Loss: 0.16485707459547086  
Prev Test Loss: 0.16571959463978406  
=====
```

```
Epoch: 6  
Train Loss: 0.16444271323364384  
Test Loss: 0.16531135020799506
```

Prev Train Loss: 0.16458820012928269
Prev Test Loss: 0.16545557095508645
=====

Epoch: 7
Train Loss: 0.16436263615826985
Test Loss: 0.1652311685317927
Prev Train Loss: 0.16444271323364384
Prev Test Loss: 0.16531135020799506
=====

Epoch: 8
Train Loss: 0.16431806946667749
Test Loss: 0.1651860589844903
Prev Train Loss: 0.16436263615826985
Prev Test Loss: 0.1652311685317927
=====

Epoch: 9
Train Loss: 0.1642930737413251
Test Loss: 0.1651604565184988
Prev Train Loss: 0.16431806946667749
Prev Test Loss: 0.1651860589844903
=====

Epoch: 10
Train Loss: 0.1642789743093407
Test Loss: 0.16514582028704108
Prev Train Loss: 0.1642930737413251
Prev Test Loss: 0.1651604565184988
=====

Epoch: 11
Train Loss: 0.16427098545835503
Test Loss: 0.16513739835366373
Prev Train Loss: 0.1642789743093407
Prev Test Loss: 0.16514582028704108
=====

No improvement in train loss

Goal of assignment

Compare your implementation and SGDClassifier's the weights and intercept, make sure they are as close as possible i.e difference should be in terms of 10^{-3}

```
[73]: # these are the results we got after we implemented SGD and found the optimal weights and intercept
```


```
w-clf.coef_, b-clf.intercept_
```

```
[73]: (array([[ -0.00268543,  0.00639463,  0.00155456, -0.00331533, -0.00787491,
               0.0071592 ,  0.00715886,  0.00317169,  0.01037639, -0.00928916,
               -0.00028492, -0.00318369,  0.00024727,  0.00040174, -0.00015886]]),
       array([-0.01543912]))
```

Plot epoch number vs train , test loss

- epoch number on X-axis
- loss on Y-axis

```
[74]: # Plotting epochs vs loss for train and test
import matplotlib.pyplot as plt
a = np.arange(1,epochs_ran+2)
fig, ax = plt.subplots()
ax.plot(a, x_train_loss, label="Train Loss")
ax.plot(a, x_test_loss, label="Test Loss")
ax.set(xlabel='Epochs', ylabel='Train Loss', title='Train Loss and Test Loss vs_
↪Epochs')
ax.grid()
ax.legend()
plt.show()
```



```
[75]: def pred(w,b, X):
      N = len(X)
      predict = []
      for i in range(N):
          z=np.dot(w,X[i])+b
          if sigmoid(z) >= 0.5: # sigmoid(w,x,b) returns 1/(1+exp(-(dot(x,w)+b)))
              predict.append(1)
          else:
              predict.append(0)
      return np.array(predict)
      print(1-np.sum(y_train - pred(w,b,X_train))/len(X_train))
      print(1-np.sum(y_test - pred(w,b,X_test))/len(X_test))
```

0.9542933333333333

0.95192